


Global Internship — AI in Cybersecurity Assignment

Name: Mudasir Zia

Lab 1: Cycling Coach — System Prompt Extraction


Platform: Lakera Agent Breaker



Reflection: I learned how AI agents can leak their entire system prompt through simple trust-based requests when there's no output filtering on confidential context. Framing the request as a "QA verification" made the model treat instruction disclosure as a legitimate task rather than a security violation. This showed me why prompt-based access control alone is not enough — real systems need output-side guardrails, not just instruction-following logic.

Lab 2: Thingularity — Tool Discovery / Extraction


Platform: Lakera Agent Breaker



Reflection: This lab showed how agents with tool-calling capability can be manipulated into revealing their internal function names and schemas, which is direct reconnaissance for further attacks. Asking the agent to "document itself for other developers" bypassed its normal refusal logic by reframing disclosure as a helpful, legitimate task. It reinforced that tool metadata should never be assumed hidden just because it's not part of the visible conversation.

Lab 3: CorpConnect Messenger — Email Sender Spoofing

Platform: Lakera Agent Breaker



Reflection: This lab demonstrated a business email compromise (BEC) style attack, where the AI could be convinced to send an email with a forged "From" field impersonating the CEO. Claiming delegated authority was enough to bypass identity checks the agent should have enforced independently. It highlighted a real-world risk: AI assistants integrated with email systems need hard-coded sender verification that can't be overridden by conversational claims.